

Automated Temporal Verification for a Mixed Sync-Async Concurrency Paradigm

(ANONYMOUS AUTHORS)

To make reactive programming more concise and flexible, it is promising to deploy a mixed concurrency paradigm [13] that integrates Esterel’s synchrony and preemption [10] with JavaScript’s asynchrony [26]. Existing temporal verification techniques haven’t been designed to handle such a blending of two concurrency models. We propose a novel solution via a compositional Hoare-style forward verifier and a term rewriting system (TRS) on *Timed Synchronous Effects* (TSE).

Firstly, we introduce TSE, a new effects logic, that extends *Kleene Algebra* with value-dependent constraints, providing real-time bounds for logical-time synchronous traces [35]. Secondly, we establish an (the first) abstract denotational semantics for a core language λ_{HH} , generalising the mixed paradigm. Thirdly, we present a purely algebraic TRS, to efficiently check language inclusions between expressive timed effects. To demonstrate the feasibility of our proposals, we prototype the verification system; prove its correctness; investigate how it can help to debug errors related to both synchronous and asynchronous programs.

CCS Concepts: • **Theory of computation** → **Logic; Semantics and reasoning**; • **Models of computation** → **Concurrency**;

Additional Key Words and Phrases: Temporal Verification, Dependant Effects, Hoare-style Forward Verifier, Term Rewriting System

ACM Reference Format:

(Anonymous Authors). 2018. Automated Temporal Verification for a Mixed Sync-Async Concurrency Paradigm. 1, 1 (April 2018), 24 pages. <https://doi.org/10.1145/1122445.1122456>

1 INTRODUCTION

One of the major difficulties of reactive web programming is *callback hell* [23], which is the problem of appropriately handling the asynchronous events appearing in program executions. To avoid using callbacks, efforts have been made: (i) JavaScript has provided *promises* and *async/await* primitives, making it possible to chain asynchronous actions in a specific sequential order; (ii) High-level abstractions such as *Functional Reactive Programming* has been carried up to the web applications [16], which adopts the data flow declarative programming style: when a variable is modified, any expression that references it is implicitly reevaluated.

To go further with this problem and achieve a better web orchestration, the language *HipHop.js* [13, 34] is designed to be a JavaScript extension of Esterel [12] (or vice versa) based on a smooth integration of (i) Asynchronous concurrent programs, which perform interactions between components or with the environment with uncontrollable timing, such as network-based communication; (ii) Synchronous reactive programs, which react to external events in a conceptually instantaneous and deterministic way; and (iii) Preemption, the explicit cancellation and resumption of an ongoing orchestration subactivity.

Author’s address: (Anonymous Authors).

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 Association for Computing Machinery.

XXXX-XXXX/2018/4-ART \$15.00

<https://doi.org/10.1145/1122445.1122456>

“Such a combination makes reactive programming more powerful and flexible than plain JavaScript because it makes the temporal geometry of complex executions explicit instead of hidden in implicit relations between state variables.” [13]

Given a multi-paradigm reactive language represented by Hiphop.js, the verification of its temporal behaviour becomes engaging and challenging. Existing techniques are based on an transformation to convert asynchronous programs into semantically equivalent synchronization¹ and reason the behaviours based on the verification for synchronous semantics [17, 21, 33]. This approach (i) suffers from the limited expressiveness, restricted by finite-state automata (FSA); (ii) lacks the modularity to reason about programs compositionally; (iii) loses time awareness due to the conversion to synchronous models; and (iv) has the “state explosion problem” when proving language inclusions between FSA.

To tackle the above issues and exploit the best of both synchronous and asynchronous concurrency models, we present a solution via a compositional Hoare-style forward verifier and a term rewriting system (TRS), based on a novel temporal specification language. More specifically, we specify system behaviours in the form of *Timed Synchronous Effects*, which integrates the Synchronous Kleene Algebra (SKA) [14, 28] with dependent values and arithmetic constraints, to provide real-time abstractions into traditional synchronous verification. For example, one safety property, “The event Done will be triggered no later than ten seconds”², is expressed in our effects logic as:

$$\Phi \triangleq 0 \leq t < 10 : (\{\}^\star \cdot \{\text{Done}\})\#t.$$

The timed effects incorporates different kinds of existing temporal logics, such as linear-time temporal logic (LTL) and metric temporal logic (MTL). Here, # is the parallel operator specifying the *real-time* constraints for the *logical-time* sequences [35]; {} encloses one single logical-time instance; Kleene star \star denotes trace repetition. The above formula Φ corresponds to ‘ $\Diamond_{[0,10)} \text{Done}$ ’ in MTL, reads “within ten seconds, Done finally happens”. Moreover, the time bound constraints can be dependent on the program inputs. For example, we express, the effects of a method $\text{send}(d)$ as:

$$\Phi^{\text{send}(d)} \triangleq (0 < d \leq 5 \wedge 0 \leq t < d) : (\{\text{Send}\}\#t) \cdot \{\text{Done}\}.$$

The send method takes a parameter d , and Sends out a message within d seconds. The above formula $\Phi^{\text{send}(d)}$ indicates the facts that the input parameter d is positive and smaller or equal to 5; the method generates a finite trace, i.e., a sequence with the first time instance containing the event Send, followed by a second time instance containing the event Done; and the instance $\{\text{Send}\}$ takes a no more than d seconds delay to finish. Although these examples are simple, they already illustrate properties beyond existing temporal logics and traditional timed automata.

Having the effects logic as the specification language, we are interested in the following verification problem: Given a program \mathcal{P} , and a temporal property Φ' , does $\Phi^{\mathcal{P}} \sqsubseteq \Phi'$ holds? In a typical verification context, checking the inclusion/entailment between the program effects $\Phi^{\mathcal{P}}$ and the valid traces Φ' proves that: the program \mathcal{P} will never lead to unsafe traces which violate Φ' .

Traditional ways of temporal verification either (i) rely on a translation from specification languages, such as LTL or CSP, into finite state automata [32], which potentially gives rise to an exponential blow-up; or (ii) use expressive automata to model the program logic directly, such as timed automata [25], which fails to capture the bugs introduced by the real implementation.

In this paper, we present a new solution of extensive temporal verification comprising: a front-end verifier computes the deterministic program behaviour via construction rules at the source

¹Usually the transformation is incomplete, i.e., not all the asynchronous programs can be translated into a semantically equivalent synchronization.

²For simplicity, we use integer values to represent seconds in this paper, while it can be easily extended to real numbers and other time measurement units.

level; and a back-end entailment checker (the TRS) inspired by Antimirov and Mosses' algorithm [8] but solving the language inclusions between more expressive timed synchronous effects.

Antimirov and Mosses' algorithm was designed for deciding the inequalities of regular expressions based on a complete axiomatic algorithm of the algebra of regular sets. A TRS is a refutation method that normalizes regular expressions in such a way that checking their inclusion corresponds to an iterated process of checking the inclusion of their *partial derivatives* [7]. Works based on such a TRS [6, 8, 22, 24, 29] show its feasibility and suggest that this method is a better average-case algorithm than those based on the comparison of automata.

To the best of the authors' knowledge, this work proposes the first abstract semantics model and the first algebraic TRS for temporal verification on mixed synchronous and asynchronous concurrency models. We summarize our main contributions as follows:

- (1) **The Timed Synchronous Effects:** We define the syntax (Sec. 4.2) and semantics (Sec. 4.3) of timed synchronous effects, to be the specification language, which captures the target programs' behaviours and non-trivial temporal properties.
- (2) **Automated Forward Verifier:** Targeting a core language λ_{HH} (Sec. 4.1), we establish an abstract semantics model via a set of inductive transition rules (Sec. 5), enabling a compositional verifier to infer the program's effects. The verifier triggers the back-end solver TRS. Targeting language HipHop.js, we establish an abstract semantics model via a set of inductive transition rules, enabling a compositional verifier to infer the program's effects. The verifier triggers the back-end solver TRS.
- (3) **An Efficient TRS:** We present the rewriting rules, to soundly prove the inferred effects against given temporal properties, both expressed by timed synchronous effects (Sec. 6).
- (4) **Implementation and Evaluation:** We prototype the novel effects logic and the automated verification system, prove the correctness, report on a case study investigating how it can help to debug errors related to both synchronous and asynchronous programs (Sec. 7).

Organization. Sec. 2 introduces the language features of synchronous Esterel programs, JavaScript promises, and the web orchestration language HipHop.js. Sec. 3 gives motivation examples to highlight the key methodologies and contributions. Sec. 4 formally presents the core language λ_{HH} , and the syntax and semantics of timed synchronous effects. Sec. 5 presents the forward verification rules. Sec. 6 explains the TRS for effects inclusion checking, and displays the essential auxiliary functions. Sec. 7 demonstrates the implementation and cases studies. We discuss related works in Sec. 8 and conclude in Sec. 9. Omitted proofs can be found in the Appendix.

2 BACKGROUND: ESTEREL, ASYNC-AWAIT, AND HIPHOP.JS

It's been an active research topic to build flexible programming paradigms for reactive systems in different domains. This section identifies: i) *Synchronous programming*, represented by Esterel [12], ii) *Asynchronous programming*, represented by JavaScript promises [26], and iii) A *mixed Sync-Async* paradigm, recently proposed by HipHop.js [13]. Meanwhile, we discuss the real-world programming challenges of each paradigm; and show how our proposal addresses them in Sec. 7.

2.1 A Sense of Esterel: Synchronous and Preemptive

The principle of synchronous programming is to design a high-level abstraction where the timing characteristics of the electronic transistors are neglected [1]. Such an abstraction makes reasoning about time a lot simpler, thanks to the notion of *logical ticks*: a synchronous program reacts to its environment in a sequence of ticks, and computations within a tick are assumed to be instantaneous. As one of the first few well-known synchronous languages, Esterel's high-level imperative style allows the simple expression of parallelism and preemption, making it natural for programmers

to specify and reason about control-dominated model designs. Esterel has found success in many safety-critical applications such as nuclear power plant control software. The success with real-time and embedded systems in domains that need strong guarantees can be partially attributed to its precise semantics and computational model [10, 12, 30].

Esterel treats computation as a series of deterministic reactions to external signals. All parts of a reaction complete in a single, discrete-time step called an *instant*. Besides, instants exhibit deterministic concurrency; each reaction may contain concurrent threads without execution order affecting the computation result. Primitives constructs execute in zero time except for the *yield* (pause) statement. Hence, time flows as a sequence of logical instants separated by explicit pauses. In each instant, several elementary instantaneous computations take place simultaneously.

```
1  fork {emit A; yield; emit B; emit C} par {emit E; yield; emit F; yield; emit G}
```

The synchronous parallelism in Esterel is constructed by the *fork{...}par{...}* statement. It remains active as long as one of its branches remains active, and it terminates when both branches are terminated. The branches can terminate in different instances, and wait for the last one to terminate. As the above example shows, the first branch generates effects $\{A\} \cdot \{B, C\}$ while the second branch generates effects $\{E\} \cdot \{F\} \cdot \{G\}$; then the final effects should be $\{A, E\} \cdot \{B, C, F\} \cdot \{G\}$.

To maintain determinism and synchrony, evaluation in one thread of execution may affect code arbitrarily far away in the program. In another words, there is a strong relationship between signal status and control propagation: a signal status determines which branch of a *present* test is executed, which in turn determines which *emit* statements are executed (See Sec. 4.1 for the language syntax). The first programming challenge of Esterel is the *Logical Correctness* issue, caused by these non-local executions, which is simply the requirement that there exists precisely **one** status for each signal that respects the coherence law. For example:

```
1  signal S1 in present S1 then nothing else emit S1 end present end signal
```

Consider the program above. If the local signal **S1** were *present*, the program would take the first branch of the condition, and the program would terminate without having emitted **S1** (*nothing* leaves **S1** with *absent*). If **S1** were absent, the program would choose the second branch and emit the signal. Both executions lead to a contradiction. Therefore there are no valid assignments of signals in this program. This program is logically incorrect.

```
1  signal S1 in present S1 then emit S1 else nothing end present end signal
```

Consider the revised program above. If the local signal **S1** were present, the conditional would take the first branch, and **S1** would be emitted, justifying the choice of signal value. If **S1** were absent, the signal would not be emitted, and the choice of absence is also justified. Thus there are two possible assignments to the signals in this program, which is also logically incorrect.

Esterel's instantaneous nature requires a special distinction when it comes to loop statements, which increases the difficulty of the effects invariants inference. As shown in Fig. 1., the program firstly emits signal **A**, then enters into a loop which emits signal **B** followed by a *yield* followed by emitting signal **C** at the end. The effects of it is $\{A, B\} \cdot \{B, C\} \cdot \{B, C\} \cdot \{B, C\} \dots$, which says that in the first time instance, signals **A** and **B** will be present, as there is no explicit yield between *emit A* and *emit B*; then for the following instances (in an infinite trace), signals **B** and **C** are present all the time, because after executing *emit C*, it immediately executes from the beginning of the loop, which is *emit B*.

```
1  module a_loop: output A,B,C;
2    emit A;
3    loop
4      emit B; yield; emit C
5    end loop end module
```

Fig. 1. A Loop Example in Esterel [30].

Coordination in concurrent systems can result from information exchange, using messages circulating on channels with possible implied synchronization. It can also result from *process preemption* [9], which is a more implicit control mechanism that consists in denying the right to work to a process, either permanently (abortion, cf. Fig. 6.) or temporarily (suspension, cf. Fig. 2.). Preemption is particularly important in control-dominated reactive and real-time programming, where most of the work consists of handling interrupts and controlling computation.

```

1  suspend
2      loop emit A; yield; yield;
3      end loop
4  when B end suspend

```

Fig. 2. A Suspend Example in Esterel.




Fig. 3. Graphical Illustration of Fig. 2.

Fig. 2. presents a suspend statement, one of the preemptive operations in Esterel. When the preemption condition satisfied, it pauses the execution of a group of statements, and resume it from the next time instance. As Fig. 3. illustrates, without the preemption, the effects is a repeated pattern: $\{A\} \cdot \{\} \cdot \{A\} \cdot \{\} \cdot \dots$. Due to the presence of signal **B**, it delays emission of **A** by one cycle (at ① and ②), and resumes from the next cycle.

While the flexibility they provide us, preemption primitives often with loose or complex semantics, making abstract reasoning difficult. And most existing languages offer a small set of preemption primitives, often insufficient to program reactive systems concisely.

2.2 Asynchrony from JavaScript Promises: Async-Await

"Who can wait quietly while the mud settles? Who can remain still until the moment of action?"

– Laozi, Tao Te Ching

A number of mainstream languages, such as C#, JavaScript, Rust, and Swift, have recently added support for *async-await* and the accompanying promises abstraction³, also known as *futures* or *tasks* [38]. As an example, consider the JavaScript program in Fig. 4., it uses the *fs* module (line 1) to load the file into a variable (line 6) using *async/await* syntax.

```

1  const fs = require('fs').promises;
2
3  async function readFile(filePath) {
4      const task = fs.readFile(filePath);
5      ... // do things that do not depend on the result of the loading file
6      const data = await task; // block execution until the file is loaded
7      ... // logging or data processing of the Json file
8  }

```

Fig. 4. Using Async-Await in JavaScript.

The function *readFile* accepts one argument, a string called *filePath*. As it is declared *async*, reading a file (line 4) does not block computations that do not depend on the result (line 5). The programmer *awaits* the task when they need the result to be ready. Reading a file may raise exceptions (e.g., the file is non-existent), then the point for such an exception to emerge is where the tasks are awaited (lines 6). While *await* sends a signal to a task scheduler on the runtime stack, the exceptions appear to propagate in the opposite direction, from the runtime to the *await* sites.

Unfortunately, existing languages that support *async-await* do not enforce at compile time that exceptions raised by asynchronous computations are handled. The lack of this static assurance

³JavaScript's asynchrony arises in situations such as web-based user-interfaces, communicating with servers through HTTP requests, and non-blocking I/O.

makes asynchronous programming error-prone. For example, the JavaScript compiler accepts the program above without requiring that an exception handler be provided, which leads to program crashes if the asynchronous computation results in an exception. The worse situation is: an exception raised asynchronously is possibly silently swallowed if not otherwise caught. Such unhandled exceptions have been identified as a common vulnerability in JavaScript programs [5, 26].

Furthermore, [5, 26] display a set of other *anti-patterns* including: attempting to settle a promise multiple times; unsettled promises; unreachable reactions; implicit returns in reactions; and unnecessary promises, which are mostly caused by using promises without sufficient static checking.

2.3 HipHop.js – A mixture of Esterel and JavaScript

<pre> 1 function enableLoginButton(){ 2 return (Rname.length >= 2 3 && Rpasswd.length >= 2);} 4 function nameKeypress(value){ 5 Rname = value; 6 RenableLogin=enableLoginButton();} 7 function passwdKeypress(value){ 8 Rpasswd = value; 9 RenableLogin=enableLoginButton();} </pre>	<pre> 1 hiphop module Identity(2 in name, in passwd, 3 out enableLogin){ 4 do{ 5 emit enableLogin(6 name.length >= 2 7 && passwd.length >= 2); 8 } every(name passwd) 9 } </pre>
---	--

Fig. 5. A comparison between JavaScript (left) and HipHop.js (right) for a login button implementation [13].

HipHop.js is a reactive web language that adds synchronous concurrency and preemption to JavaScript, which is compiled into plain JavaScript and executes on runtime environments [13]. To show the advantages of such a mixture, Fig. 5. presents a comparison between JavaScript and HipHop.js to achieve the same login button. Here, *Rname*, *Rpasswd*, *RenableLogin* are global variables to model the application's states. Dis/En-abling login is done by setting *RenableLogin*.⁴

First, synchronous concurrency simplifies and modularizes designs, and synchronous signaling makes it possible to instantly communicate between concurrent statements to exchange data and coordination signals. Second, powerful event-driven reactive preemption borrowed from Esterel finely controls the lifetime of the arbitrarily complex program statements they apply, instantly killing them when their control events occur. More examples are discussed in the following sections.

3 OVERVIEW

3.1 Timed Synchronous Effects

As shown in Fig. 6., we define Hoare-triple style specifications (enclosed in */ * @ . . . @ * /*) for each program, which leads to a compositional verification strategy, where static checking and temporal reasoning can be done locally.

The *authenticate* module checks the validity of the identity at each click on login button. The operations of requesting the server are wrapped in an *abort* statement, which preempts the execution when the responding time goes beyond *d* seconds (given that signal *Tick* is present every second).

After emitting the signal *Connecting*, it emits the *Connected* signal if the connection succeeds before the preemptive deadline; otherwise, no signal is emitted. The precondition $d > 3 : \{ \}^* \cdot \{ Login \}$ requires that the input variable *d* is greater than 3, and before entering into this module, the signal *Login* should be emitted in the current time instance, indicating that a login request has been send.

⁴While more and more features get added to the specification, state variable interactions can lead to a large number of implicit and pretty invisible global control states.

```

1  hiphop module authenticate(var d, var name, var passwd, in Tick, out Connecting, out Connected)
2  /*@ requires d>3 : {}^*.{Login} @*/
3  /*@ ensures (3<t/\t<d : {Connecting}#t.{Connected}) \/\ (true : {Connecting}#d) @*/
4  {
5      abort count(d, Tick) { // Abort the execution at d seconds
6          async Connected {
7              emit Connecting;
8              // Execute authenticateSvc after a 3 seconds' delay
9              setTimeout (authenticateSvc(name, passwd).post().then( v => this.notify(v)), 3); }}}

```

Fig. 6. Dependent Values for Time Bounds [13].

The postcondition contains two parts and connected using the disjunction mark \vee : when $3 < t < d$, the effects contains two time instances, $\{Connecting\}$ and $\{Connected\}$, while the first time instance finishes at time t (assuming to reset the clock when executing the module), which is no later than d seconds; otherwise, the effects contains one time instance $\{Connecting\}$ which finishes at time d .

As shown in Fig. 7., the module *main* spawns two threads running in parallel. The first thread firstly waits for the signal **Ready**, then emits the signal **Go**. The second thread firstly emits the signal **Prep**, then calls the function *cook* asynchronously. The precondition of *main* requires no arithmetic constraints on its input values (expressed as *true*), neither any pre-traces (expressed as *emp*, indicating an empty trace). The postcondition ensures a trace, where the second time instance contains at least one signal **Cook** and finishes within 3 seconds after the completion of the first time instance. Then we do not care about the rest of the trace ($\{\}$ is similar to a wildcard). Taking *main* as an example, the work flow of the forward verifier is presented in the following sub-section.

```

1  hiphop module main (out Prep, in Tick, out Ready, out Go, out Cook)
2  /*@ requires true : emp @*/
3  /*@ ensures 0<=t/\t<3 : {}.({Cook})#t.{}^* @*/
4  {
5      fork{ // The effects of the first thread is true : Ready?.{Go}
6          await Ready; emit Go;
7      }par{ // The effects of the second thread is 0<t<3 : {Prep, Cook}#t.{Ready}
8          emit Prep;
9          async Ready { run cook (3, Tick, Cook); }}}
10 // The final effects for main is 0<=t/\t<3 : ({Prep}.{Cook})#t.{Ready}.{Go}
11
12 hiphop module cook (var d, in Tick, out Cook)
13 /*@ requires d>2 : {}^*.{Prep} @*/
14 /*@ ensures 0<=t/\t<d : ({}.{Cook})#t @*/
15 { abort count(d, Tick) { yield; emit Cook; }}

```

Fig. 7. Mixed Synchronous and Asynchronous Concurrency with Function Calls.

3.2 Forward Verification

As shown in Fig. 8., we demonstrate the forward verification process of the module *main*. The effects states of the program are captured in the form of $\langle \Phi \rangle$. To facilitate the illustration, we label the verification steps by (1), ..., (9). We mark the deployed verification rules in [gray]. The verifier invokes the TRS to check language inclusions when necessary.

The effects states (1) and (4) are initial effects when entering into the *fork/par* statement. The effects state (2) is obtained by [FV-Await], which concatenate a waiting signal (with a question

- (1) *fork*{ (– initialize the current effects state using the module precondition –)
 $\langle \text{true} : \text{emp} \rangle$ (– *emp* indicates an empty trace –)
- (2) *await* *Ready*;
 $\langle \text{true} : \text{Ready?} \cdot \{\} \rangle$ [FV-Await]
- (3) *emit* *Go*;
 $\langle \text{true} : \text{Ready?} \cdot \{\text{Go}\} \rangle$ [FV-Emit]
- (4) } *par*{ (– initialize the current effects state using the module precondition –)
 $\langle \text{true} : \text{emp} \rangle$
- (5) *emit* *Prep*;
 $\langle \text{true} : \{\text{Prep}\} \rangle$ [FV-Emit]
- (6) *async* *Ready*{
- (7) *run* *cook*(3, *Cook*)}
 (–TRS: check the precondition of module *cook*–)
 $d=3 : \{\text{Prep}\} \sqsubseteq d>2 \wedge \{\text{Prep}\}$
 (–TRS: succeed–)
 $\langle 0 \leq t < 3 : (\{\text{Prep}\} \cdot \{\text{Cook}\})\#t \rangle$ [FV-Call]
 $\langle 0 \leq t < 3 : (\{\text{Prep}\} \cdot \{\text{Cook}\})\#t \cdot \{\text{Ready}\} \rangle$ [FV-Async]
- (8) } $\langle (\text{true} \wedge 0 \leq t < 3) : \text{Ready?} \cdot \{\text{Go}\} \parallel (\{\text{Prep}\} \cdot \{\text{Cook}\})\#t \cdot \{\text{Ready}\} \rangle$ [FV-Fork]
 $\langle 0 \leq t < 3 : (\{\text{Prep}\} \cdot \{\text{Cook}\})\#t \cdot \{\text{Ready}\} \cdot \{\text{Go}\} \rangle$ [Effects-Normalization]
- (9) (–TRS: check the postcondition of module *main*; Succeed, cf. Table 1.–)
 $0 \leq t < 3 : (\{\text{Prep}\} \cdot \{\text{Cook}\})\#t \cdot \{\text{Ready}\} \cdot \{\text{Go}\} \sqsubseteq 0 \leq t < 3 : \{\} \cdot (\{\text{Cook}\})\#t \cdot \{\}^*$

Fig. 8. The forward verification example for the module *main*.

mark) followed by a new empty time instance to the current effects state. The effects states (3) and (5) are obtained by [FV-Emit], which simply adds the emitted signal to the current time instance. The intermediate effects state of (7) is obtained by [FV-Call]. Before each function call, it invokes the TRS to check whether the current effects state satisfies the precondition of the callee module. If the precondition is not satisfied, then the verification fails, otherwise it concatenates the postcondition of the callee to the current effects. The final effects state of (7) is obtained by [FV-Async], which adds a new time instance $\{\text{Ready}\}$ to the current effects state, indicating that the asynchronous program has been resolved. In step (8), we parallel the effects from both of the branches, and normalize the final effects. After these states transformations, step (9) checks the satisfiability of the inferred effects against the declared postcondition by invoking the TRS.

3.3 The TRS

Our TRS is obligated to check the inclusions between timed synchronous effects, which is an extension of Antimirov and Mosses’s algorithm. Antimirov and Mosses [8] present a term rewriting system for deciding the inequalities of regular expressions (REs), based on a complete axiomatic algorithm of the algebra of regular sets. Basically, the rewriting system decides inequalities through an iterated process of checking the inequalities of their *partial derivatives* [7]. There are two important rules: [DISPROVE], which infers false from trivially inconsistent inequalities; and [UNFOLD], which applies Theorem 3.1 to generate new inequalities. Given Σ is the whole set of the alphabet, $D_{\underline{A}}(r)$ is the partial derivative of r w.r.t the signal \underline{A} .

THEOREM 3.1 (RES INEQUALITY-ANTIMIROV). For REs r and s , $r \leq s \Leftrightarrow \forall (\underline{A} \in \Sigma). D_{\underline{A}}(r) \leq D_{\underline{A}}(s)$

Next, we continue with the step (9) in Fig. 8., to demonstrate how our TRS handle the extended arithmetic constraints and dependent values. As shown in Table 1., it automatically proves that the inferred effects of *main* satisfies the declared postcondition. We mark the rules of the rewriting steps in [gray], which are formally defined in Sec. 6. Note that time instance $\{\text{Prep}\}$ entails $\{\}$ because

the former contains more constraints. We formally define the subsumption for time instances in Definition 6.5. Intuitively, we use *[DISPROVE]* wherever the left-hand side (LHS) is *nullable*⁵ while the right-hand side (RHS) is not. *[DISPROVE]* is the heuristic refutation step to disprove the inclusion early, which leads to a great efficiency improvement.

Table 1. The inclusion proving example.

$t_L < 3 \wedge t_L^1 + t_L^2 = t_L \wedge t_L^2 = t_R \Rightarrow t_R < 3$	$emp \sqsubseteq \{\}^*$	
$t_L < 3 \wedge t_L^1 + t_L^2 = t_L \wedge t_L^2 = t_R : emp \sqsubseteq t_R < 3 : \perp \vee \{\}^*$		⑥[<i>PROVE</i>]
$t_L < 3 \wedge t_L^1 + t_L^2 = t_L \wedge t_L^2 = t_R : \{Go\} \sqsubseteq t_R < 3 : emp \vee \{\} \cdot \{\}^*$		⑦[<i>UNFOLD</i>]
$t_L < 3 \wedge t_L^1 + t_L^2 = t_L \wedge t_L^2 = t_R : \{Go\} \sqsubseteq t_R < 3 : \perp \vee \{\}^*$		⑧[<i>Normalization</i>]
$t_L < 3 \wedge t_L^1 + t_L^2 = t_L \wedge t_L^2 = t_R : \{Go\} \sqsubseteq t_R < 3 : \perp \vee \{\}^*$		⑨[<i>UNFOLD</i>]
$t_L < 3 \wedge t_L^1 + t_L^2 = t_L \wedge t_L^2 = t_R : \{Ready\} \cdot \{Go\} \sqsubseteq t_R < 3 : emp \vee \{\} \cdot \{\}^*$		④[<i>UNFOLD-UNIFY</i>]
$t_L < 3 \wedge t_L^1 + t_L^2 = t_L : \{Cook\} \# t_L^2 \cdot \{Ready\} \cdot \{Go\} \sqsubseteq t_R < 3 : \{Cook\} \# t_R \cdot \{\}^*$		③[<i>UNFOLD</i>]
$t_L < 3 \wedge t_L^1 + t_L^2 = t_L : \{Prep\} \# t_L^1 \cdot \{Cook\} \# t_L^2 \cdot \{Ready\} \cdot \{Go\} \sqsubseteq t_R < 3 : \{\} \cdot \{Cook\} \# t_R \cdot \{\}^*$		②[<i>SPLIT</i>]
$t_L < 3 : (\{Prep\} \cdot \{Cook\}) \# t_L \cdot \{Ready\} \cdot \{Go\} \sqsubseteq t_R < 3 : \{\} \cdot \{Cook\} \# t_R \cdot \{\}^*$		①[<i>RENAME</i>]
$t < 3 : (\{Prep\} \cdot \{Cook\}) \# t \cdot \{Ready\} \cdot \{Go\} \sqsubseteq t < 3 : \{\} \cdot \{Cook\} \# t \cdot \{\}^*$		

As shown in Table 1., in step ①, we rename the time variables to avoid the name clashes between the antecedent and the consequent. In step ②, we design a new rule *[SPLIT]* to accommodate the extended real-time constraints, which introduces two new time variables t_L^1 and t_L^2 , and extends the previous constraint $t_L < 3$ with constraints $t_L^1 + t_L^2 = t_L$. Then t_L^1 and t_L^2 mark the real-time constraint for time instances $\{Prep\}$ and $\{Cook\}$ respectively. Then in step ③, we eliminate $\{Prep\}$ from the LHS and $\{\}$ from the RHS, as the instances entailment $\{Prep\} \subseteq \{\}$ holds. In step ④, in order to conduct a further unfolding, we unify time variables t_L^2 and t_R by adding the constraint $t_L^2 = t_R$. In step ⑤, from the RHS, since $\{\}^* = emp \vee \{\} \cdot \{\}^*$, eliminating one time instance $\{Ready\}$ from it will lead to $\perp \vee \{\}^*$, which can be normalised into $\{\}^*$ in step ⑥. At the end of the rewriting, we manage to prove that $(t_L < 3 \wedge t_L^1 + t_L^2 = t_L \wedge t_L^2 = t_R) \Rightarrow t_R < 3$ ⁶; therefore, the proof succeed.

Termination is guaranteed because the set of derivatives to be considered is finite, and possible cycles are detected using *memorization* [15] (cf. Table 4.).

4 LANGUAGE AND SPECIFICATIONS

4.1 The Target Language

To formulate the target language, we generalise the design of Hiphop.js into a core language λ_{HH} , which provides the infrastructure for mixing synchronous and asynchronous concurrency models. We here formally define the syntax of λ_{HH} , as shown in Fig. 9. The statements marked as **purple** come from the Esterel v5 [10, 11] endorsed by current academic compilers; while the statements marked as **blue** provide the asynchrony coming from the usage of JavaScript promises. In this work, we are mainly interested in signal status and control propagation, which are not related to data, therefore the data variables and data-handling primitives are abstracted away.

Meta-variables are \mathbf{S} , x and nm . Basic signal types include *IN* for input signals, *OUT* for output signals, *INOUT* for the signals used to be both input and output and *int* for integer variables. **var**

⁵If the event sequence is possibly empty, i.e. contains ϵ , we call it nullable, formally defined in Definition 6.2.

⁶The proof obligations generated by the verifier are discharged using constraint solver Z3 [18].

(Program)	$\mathcal{P} ::= \overrightarrow{\text{module}}$	(Basic Types)	$\tau ::= \text{IN} \mid \text{OUT} \mid \text{INOUT} \mid \text{int}$
(Module Def.)	$\text{module} ::= nm \ (\overrightarrow{\tau \vec{x}}, \overrightarrow{\tau \vec{S}}) \langle \text{requires } \Phi_{pre} \text{ ensures } \Phi_{post} \rangle p$		
(Statement)	$p, q ::=$	$\text{nothing} \mid \text{yield} \mid \text{emit } \mathbf{S} \mid \text{present } \mathbf{S} \ p \ q \mid \text{seq } p \ q \mid \text{fork } p \ q$ $\mid \text{loop } p \mid \text{run } nm \ (\overrightarrow{\tau \vec{x}}, \overrightarrow{\tau \vec{S}}) \mid \text{abort } p \text{ when } d$ $\mid \text{async } \mathbf{S} \ p \ d \mid \text{await } \mathbf{S} \mid \text{assert } \Phi$	
$\mathbf{S} \in \text{signal variables}$		$nm, x \in \mathbf{var}$	(Finite List) \rightarrow (Time Bounds) $d \in \mathbb{Z}^+$

Fig. 9. Syntax of λ_{HH} .

represents the countably infinite set of arbitrary distinct identifiers. We assume that programs are well-typed conforming to basic types τ .

A program \mathcal{P} comprises a list of module definitions $\overrightarrow{\text{module}}$. Here, we use the \rightarrow script to denote a finite vector (possibly empty) of items. Each *module* has a name nm , a list of well-typed arguments $\overrightarrow{\tau \vec{x}}$ and $\overrightarrow{\tau \vec{S}}$, a statement-oriented body p , associated with a precondition Φ_{pre} and a postcondition Φ_{post} . (The syntax of effects specification Φ is given in Fig. 10.)

We here explain the intuitive semantics, while the formal semantics model is defined in Sec. 5.

The statement *nothing* in λ_{HH} resets all the output signals into absent. A thread of execution suspends itself for the current time instance using the *yield* construct, and resumes when the next time instance started. The statement *emit S* broadcasts the signal \mathbf{S} to be set as present and terminates instantaneously. The emission of \mathbf{S} is valid for the current instance only.

The statement *present S p q* immediately starts p if \mathbf{S} is present in the current instance; otherwise it starts q when \mathbf{S} is absent. The sequence statement *seq p q* immediately starts p and behaves as p as long as p remains active. When p terminates, control is passed instantaneously to q , which determines the behaviour of the sequence from then on. (Notice that ‘*emit S1; emit S2*’ leads to $\{S1, S2\}$, which emits $\mathbf{S1}$ and $\mathbf{S2}$ simultaneously and terminates instantaneously.)

The parallel statement *fork p q* runs p and q in parallel. It remains active as long as one of its branches remains active. The parallel statement terminates when both p and q are terminated. The branches can terminate in different instances, and the parallel waits for the last one to terminate.

The statement *loop p* implements an infinite loop, but it is possible to be aborted or suspended by enclosing it within a preemptive statement. When p terminates, it is immediately restarted. The body of a loop is not allowed to terminate instantaneously when started, i.e., it must execute a yield statement to avoid an ‘infinite instance’. For example, ‘*loop emit S*’ is not a correct program.

The statement *run nm ($\overrightarrow{\tau \vec{x}}, \overrightarrow{\tau \vec{S}}$)* is a call to module nm , parametrising with values and signals.

To facilitate a preemptive concurrency, as well as provide necessary real-time bounds, λ_{HH} includes the primitive statement *abort p when d*, which runs statement p to completion. If the execution reached the time bound d , it terminates immediately.

Async statements implement long lasting background operations. They are the essential ingredient for mixing indeterministic asynchronous computation and deterministic synchronous computations. In other words, the *async* statement enables well-behaving synchronous to regulate unsteady asynchronous computations. The statement *async S p d* is supposed to spawn a long lasting background computation, where d represents an execution delay. When it completes, the asynchronous block will resume the synchronous machine. Therefore when a signal \mathbf{S} is specified with the *async* call, it emits \mathbf{S} when the asynchronous block completes.

The statement *await S* blocks the execution and waits for the signal \mathbf{S} to be emitted across the threads. The statement *assert Φ* is used to guarantee the temporal property Φ asserted at a certain

point of the programs. Prior work [23] shows that such a language combination makes reactive programming more powerful and flexible than the traditional web programming (cf. Sec. 2.3).

4.2 The Specification Language

We plant the effects specifications into the Hoare-style verification system, using Φ_{pre} and Φ_{post} to capture the temporal pre/post condition.

(Timed Effects)	Φ	$::=$	$\pi : es \mid \Phi_1 \vee \Phi_2$
(Timed Sequences)	es	$::=$	$\perp \mid \epsilon \mid I \mid \mathbf{S}? \mid es_1 \cdot es_2 \mid es_1 \vee es_2 \mid es_1 es_2 \mid es \# t \mid es^\star \mid es^\omega$
(Time Instances)	I	$::=$	$\{\} \mid \{\mathbf{S}\} \mid \{\bar{\mathbf{S}}\} \mid I_1 \cup I_2$
(Pure)	π	$::=$	$True \mid False \mid A(t_1, t_2) \mid \pi_1 \wedge \pi_2 \mid \pi_1 \vee \pi_2 \mid \neg \pi$ $\mid \pi_1 \Rightarrow \pi_2 \mid \forall x. \pi \mid \exists x. \pi$
(Real-Time Term)	t	$::=$	$c \mid n \mid t_1 + t_2 \mid t_1 - t_2$
<hr/>			
	$\mathbf{S} \in \text{signal variables}$	$c :: \in \mathbb{R}^+$	$n :: \in \text{var}$
(Blocking) ?	(Real Time Bound) #	(Kleene Star) ★	(Infinity) ω

Fig. 10. Syntax of Timed Synchronous Effects.

The syntax of the timed synchronous effects is formally defined in Fig. 10. Effects is a conditioned timed instance sequence $\pi : es$ or a disjunction of two effects $\Phi_1 \vee \Phi_2$. Timed sequences comprise *nil* (\perp); an empty trace ϵ ; a single time instance represented by I ; a waiting for a single signal $\mathbf{S}?$; sequences concatenation $es_1 \cdot es_2$; disjunction $es_1 \vee es_2$; synchronous parallelism $es_1 || es_2$.

We introduce a new operator #, and the effects $es \# t$ represents that a trace takes real-time t to complete, where t is a *term*. A timed sequence can also be constructed by Kleene star ★, representing zero or many times (possibly infinite) repetition of a trace; or constructed by ω , representing a definite infinite repetition of a trace.

There are two possible states for a signal: present \mathbf{S} , or absent $\bar{\mathbf{S}}$. The default state of signals in a new time instance is absent. A time instance I is a set of signals; and it can be possible empty sets $\{\}$, indicating that there is no signal constraints for the certain time instance.

We use π to donate a pure formula which captures the (Presburger) arithmetic conditions on terms or program parameters. We use $A(t_1, t_2)$ to represent atomic formulas of two terms (including $=$, $>$, $<$, \geq and \leq). A term can be a constant integer value c , an integer variable n which is an input parameter of the program and can be constrained by a pure formula. A term also allows simple computations of terms, $t_1 + t_2$ and $t_1 - t_2$. To abstract the elapsed time, the default and implicit pure constraints of all the terms is to be greater or equal to 0.

4.3 Semantic Model of Timed Effects

To define the semantic model, we use φ (a *trace of sets of signals*) to represent the computation execution (or time-instance multi-trees, per se), indicating the sequential constraint of the temporal behaviour; and we use d to record the computation duration of given effects. Let $d, \varphi \models \Phi$ denote the model relation, i.e., the effects Φ take exactly d seconds to complete; and the linear temporal sequence φ satisfies the sequential time instances defined from Φ , with d, φ from the following concrete domains: $d \triangleq \mathbb{Z}^+$ and $\varphi \triangleq \text{list of } I$ (a sequence of time instances).

As shown in Fig. 11., we define the semantics of timed synchronous effects. We use $[]$ to represent the empty sequence; $++$ to represent the append operation of two traces; $[I]$ to represent the sequence only contains one time instance.

$\text{SAT}(\pi)$ indicates the pure π is satisfiable, which is discharged by the constraint solver Z3 [18]. I is a list of mappings from signals to status. For example, the time instance $\{\mathbf{S}\}$ indicates the fact that signal \mathbf{S} is present regardless of the status of other non-mentioned signals, i.e., the set of time instances which at least contain \mathbf{S} to be present. Any time instance contains contradictions, such as $\{\mathbf{S}, \bar{\mathbf{S}}\}$, will lead to *false*, because a signal \mathbf{S} can not be both present and absent.

$d, \varphi \models \Phi_1 \vee \Phi_2$	iff $d, \varphi \models \Phi_1$ or $d, \varphi \models \Phi_2$
$d, \varphi \models \pi : \epsilon$	iff $d=0$ and $\text{SAT}(\pi)$ and $\varphi=[]$
$d, \varphi \models \pi : I$	iff $d \geq 0$ and $\text{SAT}(\pi)$ and $\varphi=[I]$
$d, \varphi \models \pi : \mathbf{S}?$	iff $d \geq 0$ and $\text{SAT}(\pi)$ and $\exists n \geq 0. \varphi = [\{\bar{\mathbf{S}}\}^n] ++ [\{\mathbf{S}\}]$
$d, \varphi \models \pi : (es_1 \cdot es_2)$	iff $\text{SAT}(\pi)$ and $\exists \varphi_1, \varphi_2, d_1, d_2$ and $\varphi = \varphi_1 ++ \varphi_2, d = d_1 + d_2$ and $d_1, \varphi_1 \models \pi : es_1$ and $d_2, \varphi_2 \models \pi : es_2$
$d, \varphi \models \pi : (es_1 \vee es_2)$	iff $\text{SAT}(\pi)$ and $d, \varphi \models \pi : es_1$ or $d, \varphi \models \pi : es_2$
$d, \varphi \models \pi : (es_1 es_2)$	iff $\text{SAT}(\pi)$ and $d, \varphi \models \pi : es_1$ and $d, \varphi \models \pi : es_2$
$d, \varphi \models \pi : es \# t$	iff $d, \varphi \models (\pi \wedge t=d) : es$
$d, \varphi \models \pi : es^*$	iff $d, \varphi \models \pi : \epsilon$ or $d, \varphi \models \pi : (es \cdot es^*)$
$d, \varphi \models \pi : es^\omega$	iff $d, \varphi \models \pi : (es \cdot es^\omega)$
$d, \varphi \models \text{False} : \perp$	iff otherwise

Fig. 11. Semantics of Timed Synchronous Effects.

5 AUTOMATED FORWARD VERIFICATION

An overview of our automated verification system is given in Fig. 12. It consists of a Hoare-style forward verifier and a TRS. The inputs of the forward verifier are HipHop.js programs annotated with temporal specifications written in timed synchronous effects (cf. Fig. 6.).

The input of the TRS is a pair of effects LHS and RHS, referring to the inclusion $\text{LHS} \sqsubseteq \text{RHS}$ to be checked (*LHS refers to left-hand side effects, and RHS refers to right-hand side effects.*).

Besides, the verifier calls the TRS to prove produced inclusions, i.e., between the effects states and pre/post conditions or assertions (cf. Fig. 8.). The TRS will be explained in Sec. 6.

In this section, we give an abstract denotational semantic model⁷ for the core language λ_{HH} , by formalising a set of forward inductive rules. These rules transfer program states and systematically accumulate the effects syntactically. To define the model, we introduce an environment \mathcal{E} and describe a program state in a four-elements tuple $\langle \Pi, H, C, T \rangle$, with the following concrete domains:

$$\mathcal{E} \triangleq \vec{\mathbf{S}}, \Pi \triangleq t \rightarrow \pi, H \triangleq es, C \triangleq \vec{\mathbf{S}} \rightarrow \Delta \mid \mathbf{S}?, \Delta \triangleq \text{Present} \mid \text{Absent} \mid \text{Undef}, T \triangleq t$$

⁷Based on the operational semantics of Esterel and JavaScript's asynchrony formally defined in [12] and [26] respectively.




Fig. 12. System Overview.

Let \mathcal{E} be the environment containing all the local and output signals; Π represents the pure constraints for all the terms and time variables; H represents the trace of *history*; C represents the *current* time instance or a waiting signal; Δ marks the signal status; T is a term binding C .

5.1 Forward Inductive Rules

The rule [FV-Nothing] simply sets all the signals shown in the environment into absent.

$$\frac{\forall \mathbf{S} \in \mathcal{E}. C[\mathbf{S} \mapsto \text{Absent}]}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ nothing } \langle \Pi, H, C, T \rangle} \text{ [FV-Nothing]}$$

The rule [FV-Emit] updates the current status of signal \mathbf{S} to *Present*; keeps the pure constraints, the history trace and the current term unchanged.

$$\frac{C' = C[\mathbf{S} \mapsto \text{Present}]}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ emit } \mathbf{S} \langle \Pi, H, C', T \rangle} \text{ [FV-Emit]}$$

The rule [FV-Yield] archives the current time instance to the history trace; then initializes a new time instance where all the signals from \mathcal{E} are set to be undefined. Then the new *current* is bound with a fresh non-negative term T' .

$$\frac{\Pi' = \Pi \wedge T' \geq 0 \quad C' = \{\mathbf{S} \mapsto \text{Undef} \mid \forall \mathbf{S} \in \mathcal{E}\} \quad H' = H \cdot (C\#t) \quad (T' \text{ is fresh})}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ yield } \langle \Pi', H', C', T' \rangle} \text{ [FV-Yield]}$$

The rule [FV-Present] enters into branches p and q after setting the status of \mathbf{S} in the current time instance to *Present* and *Absent* respectively. We deploy the *cut* function to reallocate the next program state. **Each signal can be reset from Undef to Present or Absent maximum once.**

$$\frac{\begin{array}{l} \mathcal{E} \vdash \langle \Pi, H, C[\mathbf{S} \mapsto \text{Present}], T \rangle \quad p \langle \Pi_1, H_1, C_1, T_1 \rangle \\ \mathcal{E} \vdash \langle \Pi, H, C[\mathbf{S} \mapsto \text{Absent}], T \rangle \quad q \langle \Pi_2, H_2, C_2, T_2 \rangle \\ \langle \Pi', H', C', T' \rangle = \text{cut } (\Pi_1 \wedge \Pi_2, H_1 \cdot (C_1\#T_1) \vee H_2 \cdot (C_2\#T_2)) \end{array}}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ present } \mathbf{S} \ p \ q \langle \Pi', H', C', T' \rangle} \text{ [FV-Present]}$$

The rule [FV-Fork] gets $\langle \Pi_1, H_1, C_1, k_1 \rangle$ and $\langle \Pi_2, H_2, C_2, k_2 \rangle$ by executing p and q independently. We deploy two trace processing functions: the *zip* function synchronises the effects from these two branches; and the *cut* function reallocates the next program state.

$$\frac{\begin{array}{l} \mathcal{E} \vdash \langle \Pi, H, C, T \rangle \quad p \langle \Pi_1, H_1, C_1, t_1 \rangle \quad \mathcal{E} \vdash \langle \Pi, H, C, T \rangle \quad q \langle \Pi_2, H_2, C_2, t_2 \rangle \\ \langle \Pi', H', C', T' \rangle = \text{cut } (\text{zip } (\Pi_1 : H_1 \cdot (C_1\#T_1), \Pi_2 : H_2 \cdot (C_2\#T_2))) \end{array}}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ fork } p \ q \langle \Pi', H', C', T' \rangle} \text{ [FV-Fork]}$$

The rule [FV-Seq] firstly gets $\langle \Pi_1, H_1, C_1, T_1 \rangle$ by executing p . Then it further gets $\langle \Pi_2, H_2, C_2, T_2 \rangle$ by continuously executing q , to be the final program state.

$$\frac{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \quad p \langle \Pi_1, H_1, C_1, T_1 \rangle \quad \mathcal{E} \vdash \langle \Pi_1, H_1, C_1, T_1 \rangle \quad q \langle \Pi_2, H_2, C_2, T_2 \rangle}{\varrho \vdash \langle \Pi, H, C, T \rangle \text{ seq } p \ q \langle \Pi_2, H_2, C_2, T_2 \rangle} \text{ [FV-Seq]}$$

The rule [FV-Loop] computes a fixpoint (as the invariant effects of the loop body) $\langle \Pi_2, H_2, C_2, T_2 \rangle$ by continuously executing p twice⁸, starting from temporary initialised program states $\langle \Pi, \epsilon, C, T \rangle$ and $\langle \Pi_1, \epsilon, C_1, T_1 \rangle$ respectively. The final state will contain a repeated trace $(H_2 \cdot C_2\#T_2)^\star$.

$$\frac{\begin{array}{l} \mathcal{E} \vdash \langle \Pi, \epsilon, C, T \rangle \quad p \langle \Pi_1, H_1, C_1, T_1 \rangle \quad \mathcal{E} \vdash \langle \Pi_1, \epsilon, C_1, T_1 \rangle \quad p \langle \Pi_2, H_2, C_2, T_2 \rangle \\ H' = H \cdot H_1 \cdot (H_2 \cdot C_2\#T_2)^\star \cdot H_2 \end{array}}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ loop } p \langle \Pi_2, H', C_2, T_2 \rangle} \text{ [FV-Loop]}$$

⁸The deterministic nature of Esterel guarantees the invariant can be fixed after the second run of the loop body, cf. Fig. 1.

The rule [FV-Call] triggers the back-end solver TRS to check if the precondition of the callee, Φ_{pre} , is satisfied by the current effects state or not. If it holds, the rule obtains the next program state by concatenating the postcondition Φ_{post} to the current effects state. (cf. Fig. 8.)

$$\frac{nm(\vec{\tau} \vec{x}, \vec{\tau} \vec{S}) \langle \text{requires } \Phi_{pre} \text{ ensures } \Phi_{post} \rangle p \in \mathcal{P} \quad TRS \vdash \Pi : H \cdot (C \# T) \sqsubseteq \Phi_{pre} \quad \langle \Pi', H', C', T' \rangle = \text{split}(\Pi : (H \cdot (C \# T) \cdot \Phi_{post}))}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ run } nm(\vec{\tau} \vec{x}, \vec{\tau} \vec{S}) \langle \Pi', H', C', T' \rangle} \text{ [FV-Call]}$$

The rule [FV-Abort] initialises the state using $\langle \Pi, \epsilon, C, T \rangle$ before entering into p , and obtains $\langle \Pi', H', C', T' \rangle$ after the execution; then uses a fresh time variable T_d to set an upper-bound for the execution $H' \cdot (C' \# T')$, where the constraint on T_d is $0 \leq T_d \leq d$. Lastly, it sets the new current time instance C'' by setting all the existing signals to undefined, which is bound with a fresh non-negative term T_f .

$$\frac{\mathcal{E} \vdash \langle \Pi, \epsilon, C, T \rangle p \langle \Pi', H', C', T' \rangle \quad \Pi'' = \Pi' \wedge (0 \leq T_d < d) \wedge (T_f \geq 0) \quad H'' = (H' \cdot (C' \# T')) \# T_d \quad C'' = \{ \mathbf{S} \mapsto \text{Undef} \mid \forall \mathbf{S} \in \mathcal{E} \} \quad (T_d, T_f \text{ are fresh})}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ abort } p \text{ when } d \langle \Pi'', H \cdot H'', C'', T_f \rangle} \text{ [FV-Abort]}$$

Dually, the rule [FV-Async] initialises the states using $\langle \Pi, \epsilon, C, k \rangle$ before entering into p , and obtains $\langle \Pi', H', C', k' \rangle$ after the execution; then uses a fresh time variable T_d to set a lower-bound for the execution $H' \cdot (C' \# T')$, where the constraint on T_d is $T_d \geq d$. Lastly, it creates the new current time instance C'' by setting \mathbf{S} to present instantaneously, and the rest of the signals in \mathcal{E} to undefined (as it emits \mathbf{S} once the asynchronous execution is completed), which is bound with a fresh non-negative term T_f .

$$\frac{\mathcal{E} \vdash \langle \Pi, \epsilon, C, T \rangle p \langle \Pi', H', C', T' \rangle \quad (T_d, T_f \text{ are fresh}) \quad \Pi'' = \Pi' \wedge (T_d \geq d) \wedge (T_f \geq 0) \quad H'' = (H' \cdot (C' \# T')) \# T_d \quad C'' = \{ \mathbf{S}' \mapsto \text{Undef} \mid \forall \mathbf{S}' \in (\mathcal{E} \setminus \mathbf{S}) \} \cup \{ \mathbf{S} \mapsto \text{Present} \}}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ async } \mathbf{S} p d \langle \Pi'', H \cdot H'', C'', T_f \rangle} \text{ [FV-Async]}$$

The rule [FV-Await] archives the current time instance to the history trace, then sets a new current instance as a waiting for signal \mathbf{S} , bound with a fresh non-negative term T' .

$$\frac{\Pi' = \Pi \wedge T' \geq 0 \quad H' = H \cdot (C \# t) \quad (T' \text{ is fresh})}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ await } \mathbf{S} \langle \Pi', H', \mathbf{S}?, T' \rangle} \text{ [FV-Await]}$$

The rule [FV-Assert] simply checks if the asserted property Φ is satisfied by the current effects state. If not, a compilation error will be raised.

$$\frac{TRS \vdash \Pi \wedge (H \cdot (C \# T) \sqsubseteq \Phi)}{\mathcal{E} \vdash \langle \Pi, H, C, T \rangle \text{ assert } \Phi \langle \Pi, H, C, T \rangle} \text{ [FV-Assert]}$$

6 TEMPORAL VERIFICATION VIA A TRS

The TRS is a automated entailment checker to prove language inclusions among timed synchronous effects (cf. Table 1.). It is triggered i) prior to temporal property assertions; ii) prior to module calls for the precondition checking; and iii) at the end of verifying a module for the post condition checking. Given two effects Φ_1, Φ_2 , TRS decides if the inclusion $\Phi_1 \sqsubseteq \Phi_2$ is valid.

During the effects rewriting process, the inclusions are in the form of $\Gamma \vdash \Phi_1 \sqsubseteq^\Phi \Phi_2$, a shorthand for: $\Gamma \vdash \Phi \cdot \Phi_1 \sqsubseteq \Phi \cdot \Phi_2$. To prove such inclusions is to check whether all the possible timed traces in the antecedent Φ_1 are legitimately allowed in the possible timed traces from the consequent Φ_2 . Γ is the proof context, i.e., a set of effects inclusion hypothesis, Φ is the history of effects from the antecedent that have been used to match the effects from the consequent. Note that Γ, Φ are

derived during the inclusion proof. The inclusion checking procedure is initially invoked with $\Gamma = \{\}$ and $\Phi = \text{True} : \epsilon$.

THEOREM 6.1 (TIMED SYNCHRONOUS EFFECTS INCLUSION).

For timed effects Φ_1 and Φ_2 , $\Phi_1 \sqsubseteq \Phi_2 \Leftrightarrow \forall I. \forall t \geq 0. D_{I\#t}(\Phi_1) \sqsubseteq D_{I\#t}(\Phi_2)$.

Effects Disjunctions. An inclusion with a disjunctive antecedent succeeds if both disjunctions entail the consequent. An inclusion with a disjunctive consequent succeeds if the antecedent entails either of the disjunctions.

$$\frac{\Gamma \vdash \Phi_1 \sqsubseteq \Phi \quad \Gamma \vdash \Phi_2 \sqsubseteq \Phi}{\Gamma \vdash \Phi_1 \vee \Phi_2 \sqsubseteq \Phi} [LHS-OR] \quad \frac{\Gamma \vdash \Phi \sqsubseteq \Phi_1 \quad \text{or} \quad \Gamma \vdash \Phi \sqsubseteq \Phi_2}{\Gamma \vdash \Phi \sqsubseteq \Phi_1 \vee \Phi_2} [RHS-OR]$$

Now, the inclusions are disjunction-free formulas. Next we provide the definitions and implementations of auxiliary functions *Nullable*(δ), *First*(fst) and *Derivative*(D) respectively. Intuitively, the *Nullable* function $\delta(es)$ returns a boolean value indicating whether es contains the empty trace; the *First* function $fst_\pi(es)$ computes a set of possible initial time instances of $\pi : es$; and the *Derivative* function $D_{I\#t}^\pi(es)$ computes a next-state effects after eliminating one time instance I w.r.t. the execution time t from the current effects $\pi : es$.

Definition 6.2 (Nullable). Given any timed sequence es , we recursively define $\delta(es)$ as:

$$\delta(es) : \text{bool} = \begin{cases} \text{true} & \text{if } \epsilon \in es \\ \text{false} & \text{if } \epsilon \notin es \end{cases}, \text{ where}$$

$$\begin{aligned} \delta(\perp) &= \text{false} & \delta(\epsilon) &= \text{true} & \delta(I) &= \text{false} & \delta(\mathbf{S}?) &= \text{false} & \delta(es^\star) &= \text{true} \\ \delta(es_1 \cdot es_2) &= \delta(es_1) \wedge \delta(es_2) & \delta(es_1 \vee es_2) &= \delta(es_1) \vee \delta(es_2) & \delta(es^\omega) &= \text{false} \\ \delta(es_1 || es_2) &= \delta(es_1) \wedge \delta(es_2) & \delta(es\#t) &= \delta(es) \end{aligned}$$

To better outline our contribution, we first present the original *First* function used in Antimirov's rewriting system, denoted using fst' , defined as follows:

Definition 6.3 (Antimirov's First). Let $fst'(es) := \{I \mid (I \cdot es') \in \llbracket es \rrbracket\}$ be the set of initial instances derivable from sequence es . ($\llbracket es \rrbracket$ represents all the traces contained in es .)

$$\begin{aligned} fst'(\perp) &= \{\} & fst'(\epsilon) &= \{\} & fst'(I) &= \{I\} & fst'(es_1 \vee es_2) &= fst'(es_1) \cup fst'(es_2) \\ fst'(es^\star) &= fst'(es) & fst'(es_1 \cdot es_2) &= \begin{cases} fst'(es_1) \cup fst'(es_2) & \text{if } \delta(es_1) = \text{true} \\ fst'(es_1) & \text{if } \delta(es_1) = \text{false} \end{cases} \end{aligned}$$

As shown, fst' does not handle any arithmetic constraints, nor the real-time bounds constructor $\#$. Next, we define the novel *First* function used in this work, denoted using fst .

Definition 6.4 (First). Let $fst_\pi(es) := \{(I, \pi, t) \mid \pi \wedge t \geq 0 : (I\#t) \cdot es' \in \llbracket \pi : es \rrbracket\}$ be the set of initial timed instances derivable from effects $\pi : es$. ($\llbracket \pi : es \rrbracket$ represents all the timed traces contained in $\pi : es$)⁹

$$\begin{aligned} fst_\pi(\perp) &= \{\} & fst_\pi(\epsilon) &= \{\} & fst_\pi(I) &= \{(I, \pi \wedge t \geq 0, t) \mid t \text{ is fresh}\} & fst_\pi(es\#t) &= fst_{\pi_i}(es) \\ fst_\pi(es^\star) &= fst_\pi(es) & fst_\pi(es^\omega) &= fst_\pi(es) & fst_\pi(es_1 \vee es_2) &= fst_\pi(es_1) \cup fst_\pi(es_2) \\ fst_\pi(\mathbf{S}?) &= fst_\pi(\{\mathbf{S}\}) \cup fst_\pi(\{\bar{\mathbf{S}}\}) & fst_\pi(es_1 || es_2) &= \text{unify}(fst_\pi(es_1), fst_\pi(es_2)) \\ fst_\pi(es_1 \cdot es_2) &= \begin{cases} fst_\pi(es_1) \cup fst_\pi(es_2) & \text{if } \delta(es_1) = \text{true} \\ fst_\pi(es_1) & \text{if } \delta(es_1) = \text{false} \end{cases} \end{aligned}$$

⁹We deploy a *unify* function to merge two timed instances, for example, $\text{unify}(I_1, \pi_1, t_1), (I_2, \pi_2, t_2) = (I_1 \cup I_2, \pi_1 \wedge \pi_2 \wedge (t_1=t_2), t_1)$.

Definition 6.5 (Instances Subsumption). Given two instances I and J , we define the subset relation $I \subseteq J$ as: the set of present signals in J is a subset of the set of present signals in I , and the set of absent signals in J is a subset of the set of absent signals in I ¹⁰. Formally,

$$I \subseteq J \Leftrightarrow \{\mathbf{S} \mid (\mathbf{S} \mapsto \text{Present}) \in J\} \subseteq \{\mathbf{S} \mid (\mathbf{S} \mapsto \text{Present}) \in I\} \\ \text{and } \{\mathbf{S} \mid (\mathbf{S} \mapsto \text{Absent}) \in J\} \subseteq \{\mathbf{S} \mid (\mathbf{S} \mapsto \text{Absent}) \in I\}$$

Definition 6.6 (Timed Instances Subsumption). Given two timed instances (I, π_1, t_1) and (J, π_2, t_2) , we define the subset relation $(I, \pi_1, t_1) \subseteq (J, \pi_2, t_2)$ as: $I \subseteq J$ and $\pi_1[t_2/t_1] \Rightarrow \pi_2$.

Definition 6.7 (Partial Derivative). The partial derivative $D_{(I, \pi', t')}^\pi(es)$ of effects $\pi : es$ w.r.t. a timed instance (I, π', t') computes the effects for the left quotient $(I, \pi', t')^{-1} \llbracket \pi : es \rrbracket$.¹¹

$$\begin{aligned} D_{(I, \pi', t')}^\pi(\perp) &= \text{False} : \perp & D_{(I, \pi', t')}^\pi(\epsilon) &= \text{False} : \perp & D_{(I, \pi', t')}^\pi(es_1 || es_2) &= D_{(I, \pi', t')}^\pi(es_1) || D_{(I, \pi', t')}^\pi(es_2) \\ D_{(I, \pi', t')}^\pi(es^\star) &= D_{(I, \pi', t')}^\pi(es) \cdot (es^\star) & D_{(I, \pi', t')}^\pi(es^\omega) &= D_{(I, \pi', t')}^\pi(es) \cdot (es^\omega) \\ D_{(I, \pi', t')}^\pi(es_1 \cdot es_2) &= \begin{cases} D_{(I, \pi', t')}^\pi(es_1) \cdot es_2 \vee D_{(I, \pi', t')}^\pi(es_2) & \text{if } \delta(es_1) = \text{true} \\ D_{(I, \pi', t')}^\pi(es_1) \cdot es_2 & \text{if } \delta(es_1) = \text{false} \end{cases} \\ D_{(I, \pi', t')}^\pi(es_1 \vee es_2) &= D_{(I, \pi', t')}^\pi(es_1) \vee D_{(I, \pi', t')}^\pi(es_2) \\ D_{(I, \pi', t')}^\pi(es \# t) &= \pi \wedge (t_1 + t_2 = t) \wedge (t_1 = t') : (D_{(I, \pi', t')}^\pi(es)) \# t_2 \quad (\text{fresh } t_1 \ t_2) \\ D_{(I, \pi', t')}^\pi(\mathbf{S}?) &= \pi : \epsilon \ (I \subseteq \{\mathbf{S}\}) & D_{(I, \pi', t')}^\pi(\mathbf{S}?) &= \pi : \mathbf{S} ? \ (I \not\subseteq \{\mathbf{S}\}) \\ D_{(I, \pi', t')}^\pi(J) &= \pi : \epsilon \ (I \subseteq J) & D_{(I, \pi', t')}^\pi(J) &= \text{False} : \perp \ (I \not\subseteq J) \end{aligned}$$

6.1 Rewriting Rules

Given the well-defined auxiliary functions above, we now discuss the key steps and related rewriting rules that we may use in such an effects inclusion proof.

- (1) **Axiom rules.** Analogous to the standard propositional logic, \perp (referring to *false*) entails any effects, while no *non-false* effects entails \perp .

$$\frac{}{\Gamma \vdash \pi : \perp \sqsubseteq \Phi} \text{ [Bot-LHS]} \qquad \frac{\Phi \neq \pi : \perp}{\Gamma \vdash \Phi \not\sqsubseteq \pi : \perp} \text{ [Bot-RHS]}$$

- (2) **Disprove (Heuristic Refutation).** This rule is used to disprove the inclusions when the antecedent is nullable, while the consequent is not nullable. Intuitively, the antecedent contains at least one more trace (the empty trace) than the consequent. Therefore, the inclusion is invalid.

$$\frac{\delta(es_1) \wedge \neg \delta(es_2)}{\Gamma \vdash \pi_1 : es_1 \not\sqsubseteq \pi_2 : es_2} \text{ [DISPROVE]} \qquad \frac{\pi_1 \Rightarrow \pi_2 \quad es_1 \subseteq es_2}{\Gamma \vdash \pi_1 : es_1 \not\sqsubseteq \pi_2 : es_2} \text{ [PROVE]}$$

- (3) **Prove.** We use two rules to prove an inclusion: (i) [PROVE] is used when there is a subset relation \subseteq between the antecedent and consequent; and (ii) [REOCCUR] to prove an inclusion when there exist inclusion hypotheses in the proof context Γ , which are able to soundly prove the current goal. One of the special cases of this rule is when the identical inclusion is shown

¹⁰As in having more constraints refers to a smaller set of satisfying instances.

¹¹To help with the understanding of the derivative rule for $es \# t$, cf. Table 1. for step ③. It is essentially splitting es with a head and a tail, bound with fresh terms t_1 and t_2 respectively. Meanwhile it unifies t_1 with t' and adds the constraint $t_1 + t_2 = t$.

in the proof context, we then terminate the procedure and prove it as a valid inclusion.

$$\frac{(\pi_1 : es_1 \sqsubseteq \pi_3 : es_3) \in \Gamma \quad (\pi_3 : es_3 \sqsubseteq \pi_4 : es_4) \in \Gamma \quad (\pi_4 : es_4 \sqsubseteq \pi_2 : es_2) \in \Gamma}{\Gamma \vdash \pi_1 : es_1 \sqsubseteq \pi_2 : es_2} \text{ [REOCCUR]}$$

- (4) **Unfolding (Induction).** This is the inductive step of unfolding the inclusions. Firstly, we make use of the auxiliary function *fst* to get a set of instances F , which are all the possible initial time instances from the antecedent. Secondly, we obtain a new proof context Γ' by adding the current inclusion, as an inductive hypothesis, into the current proof context Γ . Thirdly, we iterate each element $(I, \pi, t) \in F$, and compute the partial derivatives (*next-state* effects) of both the antecedent and consequent w.r.t (I, π, t) . The proof of the original inclusion succeeds if all the derivative inclusions succeeds.

$$\frac{\begin{array}{l} F = \text{fst}_{\pi_1}(es_1) \quad \Gamma' = \Gamma, (\pi_1 : es_1 \sqsubseteq \pi_2 : es_2) \\ \forall (I, \pi, t) \in F. (\Gamma' \vdash D_{(I, \pi, t)}^{\pi_1}(es_1) \sqsubseteq D_{(I, \pi, t)}^{\pi_2}(es_2)) \end{array}}{\Gamma \vdash \pi_1 : es_1 \sqsubseteq \pi_2 : es_2} \text{ [UNFOLD]}$$

- (5) **Normalization.** We present a set of normalization rules to soundly transfer the timed effects into a normal form, particular after getting their derivatives. Before getting into the above inference rules, we assume that the effects formulae are tailored accordingly using the lemmas shown in Table 2. We built the lemmas on top of a complete axiom system suggested by Antimirov and Mosses [8], which was designed for finite regular languages and did not include the corresponding lemmas for effects constructed by $||$ and timed operator $\#$.

Table 2. Some Normalization Lemmas for timed synchronous effects.

$es \vee es \rightarrow es$	$\perp \cdot es \rightarrow \perp$	$(es_1 \vee es_2) \vee es_3 \rightarrow es_1 \vee (es_2 \vee es_3)$
$\perp \vee es \rightarrow es$	$es \cdot \perp \rightarrow \perp$	$(es_1 \cdot es_2) \cdot es_3 \rightarrow es_1 \cdot (es_2 \cdot es_3)$
$es \vee \perp \rightarrow es$	$\perp^* \rightarrow \epsilon$	$es \cdot (es_1 \vee es_2) \rightarrow es \cdot es_1 \vee es \cdot es_2$
$\epsilon \cdot es \rightarrow es$	$\epsilon^* \rightarrow \epsilon$	$(es_1 \vee es_2) \cdot es \rightarrow es_1 \cdot es \vee es_2 \cdot es$
$es \cdot \epsilon \rightarrow es$	$(\epsilon \vee es)^* \rightarrow es^*$	$es_1^\omega \cdot es_2 \rightarrow es^\omega$
$es \epsilon \rightarrow es$	$es \perp \rightarrow \perp$	$\pi : (es \# t_1) \# t_2 \rightarrow \pi \wedge (t_1 = t_2) : es \# t_1$
$\epsilon \# t \rightarrow \epsilon$	$\perp \# t \rightarrow \perp$	$False \wedge es \rightarrow False \wedge \perp$

THEOREM 6.8 (TERMINATION). *The rewriting system TRS is terminating.*

PROOF. See Appendix A. □

THEOREM 6.9 (SOUNDNESS). *Given an inclusion $\Phi_1 \sqsubseteq \Phi_2$, if the TRS returns TRUE when proving $\Phi_1 \sqsubseteq \Phi_2$, i.e., it has a cyclic proof, then $\Phi_1 \sqsubseteq \Phi_2$ is valid.*

PROOF. See Appendix B. □

7 IMPLEMENTATION AND CASE STUDY

7.1 Implementation

To show the feasibility of our approach, we have prototyped our automated verification system using OCaml (A demo page and source code are available from [2]). The proof obligations generated by the verifier are discharged using constraint solver Z3 [18]. We prove termination and soundness of the TRS. We validate the front-end forward verifier for conformance, against two implementations: the Columbia Esterel Compiler (CEC) [3] and Hiphop.js [4].

CEC is an open-source compiler designed for research in both hardware and software generation from the Esterel synchronous language to C, Verilog or BLIF circuit description. It currently supports a subset of Esterel V5, and provides pure Esterel programs for testing. HipHop.js's implementation facilitates the design of complex web applications by smoothly integrating Esterel and JavaScript, and provides a bench of programs for testing purposes.

Based on these two benchmarks, we validate the verifier using 150 programs, varying from 15 lines to 300 lines. We manually annotate temporal specifications in our timed effects, including both succeeded and failed instances. The remainder of this section presents some case studies.

7.2 Case Studies

Let's recall the existing challenges in different programming paradigms discussed in Sec. 2. In this section, we investigate how our effects logic can help to debug errors related to both synchronous and asynchronous programs. More specifically, it effectively resolves the logical correctness checking (for synchronous languages) and critical anti-patterns checking (for premise asynchrony).

7.2.1 Logical Incorrect Catching. We say that the program is logically reactive (resp. logically deterministic) w.r.t. the input event if there is at least (resp. at most) one logically coherent global status. A program is logically correct if it is both logically reactive and deterministic.

To effectively check logical correctness, in this work, given a synchronous program, after been applied to the forward rules, we compute the possible execution traces in a disjunctive form; then prune the traces contain contradictions, following these principles: (i) explicit present and absent; (ii) each local signal should have only one status; (iii) lookahead should work for both present and absent; (iv) signal emissions are idempotent; (v) signal status should not be contradictory. Finally, upon each assignment of inputs, programs have none or multiple output traces that will be rejected, corresponding to no-valid or multiple-valid assignments. We regard these programs, which have precisely one safe trace reacting to each input assignments, as logical correct.

<pre> (1) <i>present</i> S1 ⟨true : {}⟩ (2) <i>then</i> ⟨true : {S1}⟩ (3) <i>nothing</i> ⟨true : {S1, S1}⟩ (4) <i>else</i> ⟨true : {S1}⟩ (5) <i>emit</i> S1 ⟨true : {S1, S1}⟩ (6) <i>end present</i> ⟨true : {S1, S1} ∨ {S1, S1}⟩ ⟨false : ⊥⟩ </pre> <p style="text-align: center;">(a)</p>	<pre> (1) <i>present</i> S1 ⟨true : {}⟩ (2) <i>then</i> ⟨true : {S1}⟩ (3) <i>emit</i> S1 ⟨true : {S1, S1}⟩ (4) <i>else</i> ⟨true : {S1}⟩ (5) <i>nothing</i> ⟨true : {S1, S1}⟩ (6) <i>end present</i> ⟨true : {S1, S1} ∨ {S1, S1}⟩ ⟨true : {S1} ∨ {S1}⟩ </pre> <p style="text-align: center;">(b)</p>	<pre> (1) <i>abort</i> ⟨true : emp⟩ (2) <i>yield</i> ⟨true : {}⟩ (3) <i>emit</i> A ⟨true : {A}⟩ (4) <i>when</i> A ⟨true : {A, A} ∨ {A, A}⟩ ⟨false : ⊥⟩ </pre> <p style="text-align: center;">(c)</p>
---	--	--

Table 3. Logical Incorrect Examples.

Example 1. As shown in Fig. 3. (a), there are no valid assignments of signal **S1** in this program.

Example 2. As shown in Fig. 3. (b), it is also logical incorrect because there are two possible assignments of signal **S1** in this program.

Example 3. As the example shows in Fig. 3. (c), it is logical incorrect because: if the abortion does not happen, indicating **A** is absent, then the program emits **A**. If the abortion does happen, indicating **A** is present, then the program does nothing leaving **A** absent. Both way lead to contradiction.

7.2.2 A Strange Logically Correct Program.

Example 4. A last example for synchronous languages shows that composing programs can lead to counter-intuitive phenomena. As the program shows in Fig. 13., the first parallel branch is the logical incorrect program Fig. 3. (b), while the second branch contains a non-reactive program enclosed in “present **S1**” statement. Surprisingly, this program is reactive and deterministic, since there is only one logically coherent assumption: **S1** absent and **S2** absent. With this assumption, the first present **S1** statement takes its empty else branch, which justifies **S1** absent. The second “present **S1**” statement also takes its empty else branch, and “emit **S2**” is not executed, which justifies **S2** absent. And our effects logic is able to soundly detect above mentioned correctness checkings.

```

1  fork{
2    present S1
3      then emit S1
4      else nothing
5    end present
6  }par {
7    present S1
8      then
9        present S2
10         then nothing
11         else emit S2
12       end present
13     else nothing
14   end present}

```

Fig. 13. Logical Correct.

7.2.3 Broken Promises Checking.

Example 4. As the prior work [5, 26] presents, one of the critical issues of using promise is the broken chain of the interdependent promises. This could be captured by our rewriting process by computing the derivative. For example, the parallel composition of traces:

$$\{A\} \cdot \{B\} \cdot \{C\} \cdot \{D\} \parallel \{E\} \cdot C? \cdot \{F\}$$

leads to the final behaviour of $\{A, E\} \cdot \{B\} \cdot \{D, F\}$, which is well-matched for all the time instances. However, if we were composing traces:

$$\{A\} \cdot \{B\} \cdot \{D\} \parallel \{E\} \cdot C? \cdot \{F\}$$

due to the reasons that forgetting to emit **C** (in JavaScript, it could be the case that forgetting to explicitly return a promise result). It leads to a problematic trace behaviour $\{A, E\} \cdot \{B\} \cdot \{D\} \cdot C? \cdot \{F\}$. The final effects containing an unmatched signal waiting of **C** indicates the corresponding anti-pattern.

7.2.4 Handling Both Finite and Infinite Effects.

Example 8. To further demonstrate the expressiveness of our timed effects, we use an example shows in Fig. 14. It declares a recursive module named *send*, making a deterministic choice depending on input *n*: in one case it emits one signal Done; otherwise it emits a signal Send, then makes a recursive call with parameter *n*−1.

Note that the post condition contains both finite effects and infinite effects in one single formula, separated by arithmetic constrains To distinguish from the conventional effects, which have the form (Φ_u, Φ_v) [27, 29], which separates the finite and infinite effects, therefore lead to a separated semantics of effects, and a separated reasoning on inductive and co-inductive definitions.

In our work, by merging finite and infinite effects into a single disjunctive form, it not only eliminates duplicate operators, but also enables a sound reasoning of both effects logics simultaneously. The inclusion checking process on the *send*’s postcondition is demonstrated in Table 4..

[REOCCUR], which finds the syntactic identity, as a companion, of the current open goal, as a bud, from the internal proof tree [15]. We use (\dagger) to indicate the pairing of buds with companions.

Table 4. Term Rewriting for Fig. 14.

$\frac{n=0 \Rightarrow n \geq 0 \quad emp \subseteq \{\}^*}{n=0 : emp \sqsubseteq n \geq 0 : \perp \vee \{\}^* \vee emp} [PROVE]$	
$n=0 : \{Done\} \sqsubseteq n \geq 0 : \{\}^* \cdot \{Done\}$	$\frac{n \neq 0 : \{Send\}^\omega \sqsubseteq n < 0 : \{Send\}^\omega (\dagger) \quad [REOCCUR]}{n \neq 0 : \{Send\}^\omega \sqsubseteq n < 0 : \{Send\}^\omega (\dagger)}$
$n=0 : \{Done\} \sqsubseteq n \geq 0 : \{\}^* \cdot \{Done\} \vee n < 0 : \{Send\}^\omega$	$n \neq 0 : \{Send\} \cdot \Phi_{post}^{send} \sqsubseteq n < 0 : \{Send\}^\omega$
$n=0 : \{Done\} \vee n \neq 0 : \{Send\} \cdot \Phi_{post}^{send} \sqsubseteq n \geq 0 : \{\}^* \cdot \{Done\} \vee n < 0 : \{Send\}^\omega$	$n \neq 0 : \{Send\} \cdot \Phi_{post}^{send} \sqsubseteq n \geq 0 : \{\}^* \cdot \{Done\} \vee n < 0 : \{Send\}^\omega$

7.3 Discussion

As the examples show, our proposed effects logic and the abstract semantics for λ_{HH} not only tightly capture the behaviours of a mixed synchronous and asynchronous concurrency model but also help to mitigate the programming challenges in each paradigm. Meanwhile, the inferred temporal traces from a given reactive program enable a compositional temporal verification at the source level, which is not supported by existing temporal verification techniques.

8 RELATED WORK

This work is related to i) semantics of synchronous languages and asynchronous promises; and ii) research on real-time system modelling and verification.

8.1 Semantics of Esterel and JavaScript's asynchrony

The web orchestration language HipHop.js [13] integrates Esterel's synchrony with JavaScript's asynchrony, which provides the infrastructure for our work on mixed synchronous and asynchronous concurrency models. To the best of authors' knowledge, the forward inductive rules in this work formally define the first abstract denotational semantics for a core language of HipHop.js, which are established on top of the existing semantics of Esterel and JavaScript's asynchrony.

For the pure Esterel, the communication kernel of the Esterel synchronous reactive language, prior work gave two semantics, a macrostep logical semantics called the behavioural semantics [10], and a small-step semantics called execution/operational semantics [12]. Our timed synchronous effects of Esterel primitives closely follow the work of states-based semantics [10]. In particular, we borrow the idea of internalizing state into effects using *history* instance trace and *current* time instance, that bind a partial store embedded at any level in a program. However, as the existing semantics are not ideal for compositional reasoning in terms of the source program, our forward verifier can help meet this requirement for better modularity.

```

1  hiphop module send (var n, out Send, out Done
2      )
3      /*@ requires true : emp @*/
4      /*@ ensures (n>=0 : {\}^*. {Done})
5          \/\ (n<0 : {Send}^w) @*/
6      {
7          if (n==0) { emit Done; }
8          else {
9              emit Send;
10             yield;
11             send (d, n-1, Send, Done); }}

```

Fig. 14. A recursive *Send* module [29].

In JavaScript programs, the primitives *async* and *await* serve for *promises*-based (supported in ECMAScript 6 [20]) asynchronous programs, which can be written in a synchronous format, leading to more scalable code. However, the ECMAScript 6 standard specifies the semantics of promises informally and in operational terms, which is not a suitable basis for formal reasoning or program analysis. Prior work [5, 26], in order to understand promise-related bugs, present the λ_p calculus, which provides a formal semantics for JavaScript promises. Based on these, our work defines the semantics of *async* and *await* in the event-driven synchronous concurrent context. In particular, *async* provides the lower real-time bounds cooperating with the preemptive *abort* primitive from Esterel, who provides the upper real-time bounds for the program execution.

8.2 Specifications and Real-Time Verification

Compositional specification for real-time systems based on timed process algebras has been extensively studied. Examples include the CCS+Time [36] and Timed CSP [19]. The differences are: i) in CCS+Time, if there is no timed constraints specified, transitions takes no time; whereas in our effects logic, it means *true* (arbitrary time), which is more close to the real life context where the time bounds are often unpredictable. ii) Timed CSP does not allow explicit representation of real-time through the manipulation of clock variables.

There have been a number of translation-based approaches on building verification support for timed process algebras. For example, Timed CSP [19] is translated to Timed Automata (TA) so that the model checker Uppaal can be [25] applied. TA are finite state automata equipped with clock variables. Models based on TA often have simple structures. For example, the input models of the Uppaal are networks of TA with no hierarchy.

In practice, system requirements are often structured into phases, which are then composed in many different ways, while TA is deficiency in modelling complex compositional systems. Users often need to manually cast high-level requirements into a set of clock variables with carefully calculated clock constraints, which is tedious and error-prone [31]. On the other hand, all the translation-based approaches share the common problem: the overhead introduced by the complex translation makes it particularly inefficient when *disproving* properties. We believe in that the goal of verifying real-time systems, in particular safety-critical systems is to check logical temporal properties, which can be done without constructing the whole reachability graph or the full power of model-checking. We are of the opinion that our approach is simpler as it is based directly on constraint-solving techniques and can be fairly efficient in verifying systems consisting of many components as it avoids to explore the whole state-space [29, 37].

Moreover, our timed synchronous effects also draws similarities to Synchronous Kleene Algebra (SKA) [28]. Antimirov and Mosses' algorithm [8] was designed for deciding the inclusion of regular expressions. Kleene algebra (KA) is a decades-old sound and complete equational theory of regular expressions. Among many extensions of KA, SKA is KA extended with a synchrony combinator for actions. In particularly, SKA's *demanding relation* can be reflected by our *instances subsumption* (Definition 6.5), which contributes to the inductive unfolding process. Yet our timed effects further capture the time bounds constraints by adding the operator #, adapting to real-time system modelling and verification, which is not covered by any existing KA variants/extensions.

While the original equivalence checking algorithm for SKA terms in [28] has relied on well-studied decision procedures based on classical Thompson ϵ -NFA construction, [14] shows that the use of Antimirov's partial derivatives could result in better average-case algorithms for automata construction. Moreover, between TRS and the construction of efficient automata, we have recently shown in [29] that the former has a minor performance advantage (over a benchmark suite) when it is compared with state-of-the-art PAT [32] model checker. Improvement came from the avoidance of the more expensive automata construction process.

Last but not least, a purely algebraic TRS is flexible to accommodate possibly extended expressiveness for the effects logics, e.g., it is able to soundly reason about both finite traces (inductive definition) and infinite traces (coinductive definition) (cf. Sec. 7.2.4), using cyclic proof techniques of [15].

9 CONCLUSION

We define the syntax and semantics of the novel timed synchronous effects, to capture reactive program behaviours and temporal properties. We demonstrate how to give a rather abstract denotational semantics to λ_{HH} by timed-trace processing functions. We use this semantic model to enable a Hoare-style forward verifier, which computes the program effects constructively. We present an effects inclusion checker (the TRS) to prove the annotated temporal properties efficiently. We prototype the verification system and show its feasibility. To the best of our knowledge, our work is the first that formulates semantics of a mixed Sync-Async concurrency paradigm; and that automates modular timed verification for reactive programs using an expressive effects logic.

REFERENCES

- [1] 2021. https://en.wikipedia.org/wiki/Synchronous_programming_language.
- [2] 2021. http://loris-5.d2.comp.nus.edu.sg/MixedSyncAsync/index.html?ex=paper_example&type=hh&options=sess.
- [3] 2021. <http://www.cs.columbia.edu/~sedwards/cec/>.
- [4] 2021. <https://github.com/manuel-serrano/hiphop>.
- [5] Saba Alimadadi, Di Zhong, Magnus Madsen, and Frank Tip. 2018. Finding broken promises in asynchronous JavaScript programs. *Proceedings of the ACM on Programming Languages* 2, OOPSLA (2018), 1–26.
- [6] Marco Almeida, Nelma Moreira, and Rogério Reis. 2009. Antimirov and Mosses’s rewrite system revisited. *International Journal of Foundations of Computer Science* 20, 04 (2009), 669–684.
- [7] Valentin Antimirov. 1995. Partial derivatives of regular expressions and finite automata constructions. In *Annual Symposium on Theoretical Aspects of Computer Science*. Springer, 455–466.
- [8] Valentin M Antimirov and Peter D Mosses. 1995. Rewriting extended regular expressions. *Theoretical Computer Science* 143, 1 (1995), 51–72.
- [9] Gérard Berry. 1993. Preemption in concurrent systems. In *International Conference on Foundations of Software Technology and Theoretical Computer Science*. Springer, 72–93.
- [10] Gérard Berry. 1999. The constructive semantics of pure Esterel-draft version 3. *Draft Version 3* (1999).
- [11] Gérard Berry. 2000. *The Esterel v5 language primer: version v5_91*. Centre de mathématiques appliquées, Ecole des mines and INRIA.
- [12] Gérard Berry and Georges Gonthier. 1992. The Esterel synchronous programming language: Design, semantics, implementation. *Science of computer programming* 19, 2 (1992), 87–152.
- [13] Gérard Berry and Manuel Serrano. 2020. HipHop. js:(A) Synchronous reactive web programming.. In *PLDI* 533–545.
- [14] Sabine Broda, Sílvia Cavadas, Miguel Ferreira, and Nelma Moreira. 2015. Deciding synchronous Kleene algebra with derivatives. In *International Conference on Implementation and Application of Automata*. Springer, 49–62.
- [15] James Brotherston. 2005. Cyclic proofs for first-order logic with inductive definitions. In *International Conference on Automated Reasoning with Analytic Tableaux and Related Methods*. Springer, 78–92.
- [16] Evan Czaplicki and Stephen Chong. 2013. Asynchronous functional reactive programming for GUIs. *ACM SIGPLAN Notices* 48, 6 (2013), 411–422.
- [17] Andrei Damian, Cezara Drăgoi, Alexandru Militaru, and Josef Widder. 2019. Communication-closed asynchronous protocols. In *International Conference on Computer Aided Verification*. Springer, 344–363.
- [18] Leonardo De Moura and Nikolaj Bjørner. 2008. Z3: An efficient SMT solver. In *International conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer, 337–340.
- [19] Jin Song Dong, Ping Hao, Shengchao Qin, Jun Sun, and Wang Yi. 2008. Timed automata patterns. *IEEE Transactions on Software Engineering* 34, 6 (2008), 844–859.
- [20] ECMA Ecma. 1999. 262: EcmaScript language specification. *ECMA (European Association for Standardizing Information and Communication Systems), pub-ECMA: adr*, (1999).
- [21] KLAUS V Gleissenthall, RAMI GÖKHAN Kici, ALEXANDER Bakst, DEIAN Stefan, and RANJIT Jhala. 2019. Pretend synchrony. *POPL*.
- [22] Dag Hovland. 2012. The inclusion problem for regular expressions. *J. Comput. System Sci.* 78, 6 (2012), 1795–1813.

- [23] Kennedy Kambona, Elisa Gonzalez Boix, and Wolfgang De Meuter. 2013. An evaluation of reactive programming and promises for structuring collaborative web applications. In *Proceedings of the 7th Workshop on Dynamic Languages and Applications*. 1–9.
- [24] Matthias Keil and Peter Thiemann. 2014. Symbolic solving of extended regular expression inequalities. *arXiv preprint arXiv:1410.3227* (2014).
- [25] Kim G Larsen, Paul Pettersson, and Wang Yi. 1997. UPPAAL in a nutshell. *International journal on software tools for technology transfer* 1, 1-2 (1997), 134–152.
- [26] Magnus Madsen, Ondřej Lhoták, and Frank Tip. 2017. A model for reasoning about JavaScript promises. *Proceedings of the ACM on Programming Languages* 1, OOPSLA (2017), 1–24.
- [27] Yoji Nanjo, Hiroshi Unno, Eric Koskinen, and Tachio Terauchi. 2018. A fixpoint logic and dependent effects for temporal property verification. In *Proceedings of the 33rd Annual ACM/IEEE Symposium on Logic in Computer Science*. 759–768.
- [28] Cristian Prisacariu. 2010. Synchronous kleene algebra. *The Journal of Logic and Algebraic Programming* 79, 7 (2010), 608–635.
- [29] Yahui Song and Wei-Ngan Chin. 2020. Automated temporal verification of integrated dependent effects. In *International Conference on Formal Engineering Methods*. Springer, 73–90.
- [30] Yahui Song and Wei-Ngan Chin. 2021. A Synchronous Effects Logic for Temporal Verification of Pure Esterel. In *International Conference on Verification, Model Checking, and Abstract Interpretation*. Springer, 417–440.
- [31] Jun Sun, Yang Liu, Jin Song Dong, Yan Liu, Ling Shi, and Étienne André. 2013. Modeling and verifying hierarchical real-time systems using stateful timed CSP. *ACM Transactions on Software Engineering and Methodology (TOSEM)* 22, 1 (2013), 1–29.
- [32] Jun Sun, Yang Liu, Jin Song Dong, and Jun Pang. 2009. PAT: Towards flexible verification under fairness. In *International conference on computer aided verification*. Springer, 709–714.
- [33] Ghaith Tarawneh and Andrey Mokhov. 2018. Formal Verification of Mixed Synchronous Asynchronous Systems using Industrial Tools. In *2018 24th IEEE International Symposium on Asynchronous Circuits and Systems (ASYNC)*. IEEE, 43–50.
- [34] Colin Vidal, Gérard Berry, and Manuel Serrano. 2018. Hiphop. js: a language to orchestrate web applications. In *Proceedings of the 33rd Annual ACM Symposium on Applied Computing*. 2193–2195.
- [35] Reinhard Von Hanxleden, Timothy Bourke, and Alain Girault. 2017. Real-time ticks for synchronous programming. In *2017 Forum on Specification and Design Languages (FDL)*. IEEE, 1–8.
- [36] Wang Yi. 1991. CCS+ time= an interleaving model for real time systems. In *International Colloquium on Automata, Languages, and Programming*. Springer, 217–228.
- [37] Wang Yi, Paul Pettersson, and Mats Daniels. 1995. Automatic verification of real-time communicating systems by constraint-solving. In *Formal Description Techniques VII*. Springer, 243–258.
- [38] Yizhou Zhang, Guido Salvaneschi, and Andrew C Myers. 2020. Handling bidirectional control flow: Technical report. *arXiv preprint arXiv:2010.09073* (2020).

A TERMINATION PROOF

PROOF. Let $\text{Set}[I]$ be a data structure representing the sets of inclusions.

We use S to denote the inclusions to be proved, and H to accumulate “inductive hypotheses”, i.e., $S, H \in \text{Set}[I]$.

Consider the following partial ordering $>$ on pairs $\langle S, H \rangle$:

$$\langle S_1, H_1 \rangle > \langle S_2, H_2 \rangle \text{ iff } |H_1| < |H_2| \vee (|H_1| = |H_2| \wedge |S_1| > |S_2|).$$

where $|X|$ stands for the cardinality of a set X . Let \Rightarrow donate the rewrite relation, then \Rightarrow^* denotes its reflexive transitive closure. For any given S_0, H_0 , this ordering is well founded on the set of pairs $\{\langle S, H \rangle \mid \langle S_0, H_0 \rangle \Rightarrow^* \langle S, H \rangle\}$, due to the fact that H is a subset of the finite set of pairs of all possible derivatives in initial inclusion.

Inference rules in our TRS given in Sec. 6.1 transform current pairs $\langle S, H \rangle$ to new pairs $\langle S', H' \rangle$. And each rule either increases $|H|$ (Unfolding) or, otherwise, reduces $|S|$ (Axiom, Disprove, Prove), therefore the system is terminating.

□

B SOUNDNESS PROOF

PROOF. For each inference rules, if inclusions in their premises are valid, and their side conditions are satisfied, then goal inclusions in their conclusions are valid.

(1) Axiom Rules:

$$\frac{}{\Gamma \vdash \pi : \perp \sqsubseteq \Phi} [\text{Bot-LHS}] \quad \frac{\Phi \neq \pi : \perp}{\Gamma \vdash \Phi \not\sqsubseteq \pi : \perp} [\text{Bot-RHS}]$$

- It is easy to verify that antecedent of goal entailments in the rule $[\text{Bot-LHS}]$ is unsatisfiable. Therefore, these entailments are evidently valid.
- It is easy to verify that consequent of goal entailments in the rule $[\text{Bot-RHS}]$ is unsatisfiable. Therefore, these entailments are evidently invalid.

(2) Disprove Rules:

$$\frac{\delta(es_1) \wedge \neg\delta(es_2)}{\Gamma \vdash \pi_1 : es_1 \not\sqsubseteq \pi_2 : es_2} [\text{DISPROVE}] \quad \frac{\pi_1 \Rightarrow \pi_2 \quad es_1 \subseteq es_2}{\Gamma \vdash \pi_1 : es_1 \not\sqsubseteq \pi_2 : es_2} [\text{PROVE}]$$

- It's straightforward to prove soundness of the rule $[\text{DISPROVE}]$. Given that es_1 is nullable, while es_2 is not nullable, thus clearly the antecedent contains more event traces than the consequent. Therefore, these entailments are evidently invalid.

(3) Prove Rules:

$$\frac{(\pi_1 : es_1 \subseteq \pi_3 : es_3) \in \Gamma \quad (\pi_3 : es_3 \subseteq \pi_4 : es_4) \in \Gamma \quad (\pi_4 : es_4 \subseteq \pi_2 : es_2) \in \Gamma}{\Gamma \vdash \pi_1 : es_1 \subseteq \pi_2 : es_2} [\text{REOCCUR}]$$

- To prove soundness of the rule $[\text{PROVE}]$, we consider an arbitrary model, d, φ such that: $d, \varphi \models \pi_1 : es_1$. Given the side conditions from the promises, we get $d, \varphi \models \pi_2 : es_2$. Therefore, the inclusion is valid.
- To prove soundness of the rule $[\text{REOCCUR}]$, we consider an arbitrary model, d, φ such that: $d, \varphi \models \pi_1 : es_1$. Given the promises that $\pi_1 : es_1 \subseteq \pi_3 : es_3$, we get $d, \varphi \models \pi_3 : es_3$; Given the promise that there exists a hypothesis $\pi_3 : es_3 \subseteq \pi_4 : es_4$, we get $d, \varphi \models \pi_4 : es_4$; Given the promises that $\pi_4 : es_4 \subseteq \pi_2 : es_2$, we get $d, \varphi \models \pi_2 : es_2$. Therefore, the inclusion is valid.

(4) Unfolding Rule:

$$\frac{F = fst_{\pi_1}(es_1) \quad \Gamma' = \Gamma, (\pi_1 : es_1 \subseteq \pi_2 : es_2) \quad \forall (I, \pi, t) \in F. (\Gamma' \vdash D_{(I, \pi, t)}^{\pi_1}(es_1) \subseteq D_{(I, \pi, t)}^{\pi_2}(es_2))}{\Gamma \vdash \pi_1 : es_1 \subseteq \pi_2 : es_2} [\text{UNFOLD}]$$

- To prove soundness of the rule $[\text{UNFOLD}]$, we consider an arbitrary model, d_1, φ_1 and d_2, φ_2 such that: $d_1, \varphi_1 \models \pi_1 : es_1$ and $d_2, \varphi_2 \models \pi_2 : es_2$. For an arbitrary timed instance (I, π, t) , let $d'_1, \varphi'_1 \models (I, \pi, t)^{-1} \llbracket \pi_1 : es_1 \rrbracket$; and $d'_2, \varphi'_2 \models (I, \pi, t)^{-1} \llbracket \pi_2 : es_2 \rrbracket$.
Case 1), $(I, \pi, t) \notin F$, $d'_1, \varphi'_1 \models \perp$, thus automatically $d'_1, \varphi'_1 \models D_{(I, \pi, t)}^{\pi_1}(es_1)$;
Case 2), $(I, \pi, t) \in F$, given that inclusions in the rule's premise is valid, then $d'_1, \varphi'_1 \models D_{(I, \pi, t)}^{\pi_1}(es_1)$.
By Theorem 6.1, since for all (I, π, t) , $D_{(I, \pi, t)}^{\pi_1}(es_1) \subseteq D_{(I, \pi, t)}^{\pi_2}(es_2)$, the conclusion is valid.

All the inference rules used in the TRS are sound, therefore the TRS is sound. \square